

Feature Detection, Matching and its Applications

Alex Lin



Outline

- ***Introduction***
- ***What is feature***
- ***Where is feature***
- ***Feature Matching***

Building a Panorama



M. Brown and D. G. Lowe. Recognising Panoramas. ICCV 2003

How do we build a panorama?





- We need to match (align) images
- Global methods sensitive to occlusion, lighting, parallax effects. So look for local features that match well.
- How would you do it by eye?



Matching with Features




- Detect feature points in both images



Matching with Features



- Detect feature points in both images
- Find corresponding pairs



Matching with Features



- Detect feature points in both images
- Find corresponding pairs
- Use these pairs to align images (homography)



Matching with Features

- Problem 1:
 - Detect the *same* point *independently* in both images



no chance to match!

We need a repeatable detector

Matching with Features

- Problem 2:
 - For each point correctly recognize the corresponding one




We need a reliable and distinctive descriptor

More motivation...

- Feature points are used also for:
 - Image alignment (homography, fundamental matrix)
 - 3D reconstruction
 - Motion tracking
 - Object recognition
 - Indexing and database retrieval
 - Robot navigation
 - ... other

Query book




Book identification



Selecting Good Features

- What's a “good feature” ?
 - Satisfies brightness constancy—looks the same in both images
 - Has sufficient texture variation
 - Does not have too much texture variation
 - Corresponds to a “real” surface patch—see below:




- Does not deform too much over time



An introductory example:


Harris corner detector




C.Harris, M.Stephens. "A Combined Corner and Edge Detector". 1988

The Basic Idea


- We should easily localize the point by looking through a small window
- Shifting a window in *any direction* should give *a large change* in intensity




Harris Detector: Basic Idea



“flat” region:
no change as shift
window in all direc
tions



“edge”:
no change as shift
window along the
edge direction



“corner”:
significant change as
shift window in all dire
ctions

Harris Detector: Mathematics

Window-averaged change of intensity induced by shifting the image data by $[u, v]$:


$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Window
function


Shifted
intensity

Intensity

Window function $w(x, y) =$



1 in window, 0 outside



Gaussian

Taylor series approx to shifted image

$$\begin{aligned} E(u, v) &\approx \sum_{x, y} w(x, y) [I(x, y) + uI_x + vI_y - I(x, y)]^2 \\ &= \sum_{x, y} w(x, y) [uI_x + vI_y]^2 \\ &= \sum_{x, y} w(x, y) (u \quad v) \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix} \begin{pmatrix} u \\ v \end{pmatrix} \end{aligned}$$

Harris Detector: Mathematics

Expanding $I(x,y)$ in a Taylor series expansion, we have, for small shifts $[u, v]$, a *bilinear* approximation:

$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a 2×2 matrix computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

M is also called “structure tensor”


Harris Detector: Mathematics

Intensity change in shifting window: eigenvalue analysis



$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix} \quad \lambda_1, \lambda_2 - \text{eigenvalues of } M$$

Ellipse $E(u, v) = \text{const}$

Iso-intensity contour of $E(u, v)$





Selecting Good Features





λ_1 and λ_2 are large

Selecting Good Features



large λ_1 , small λ_2

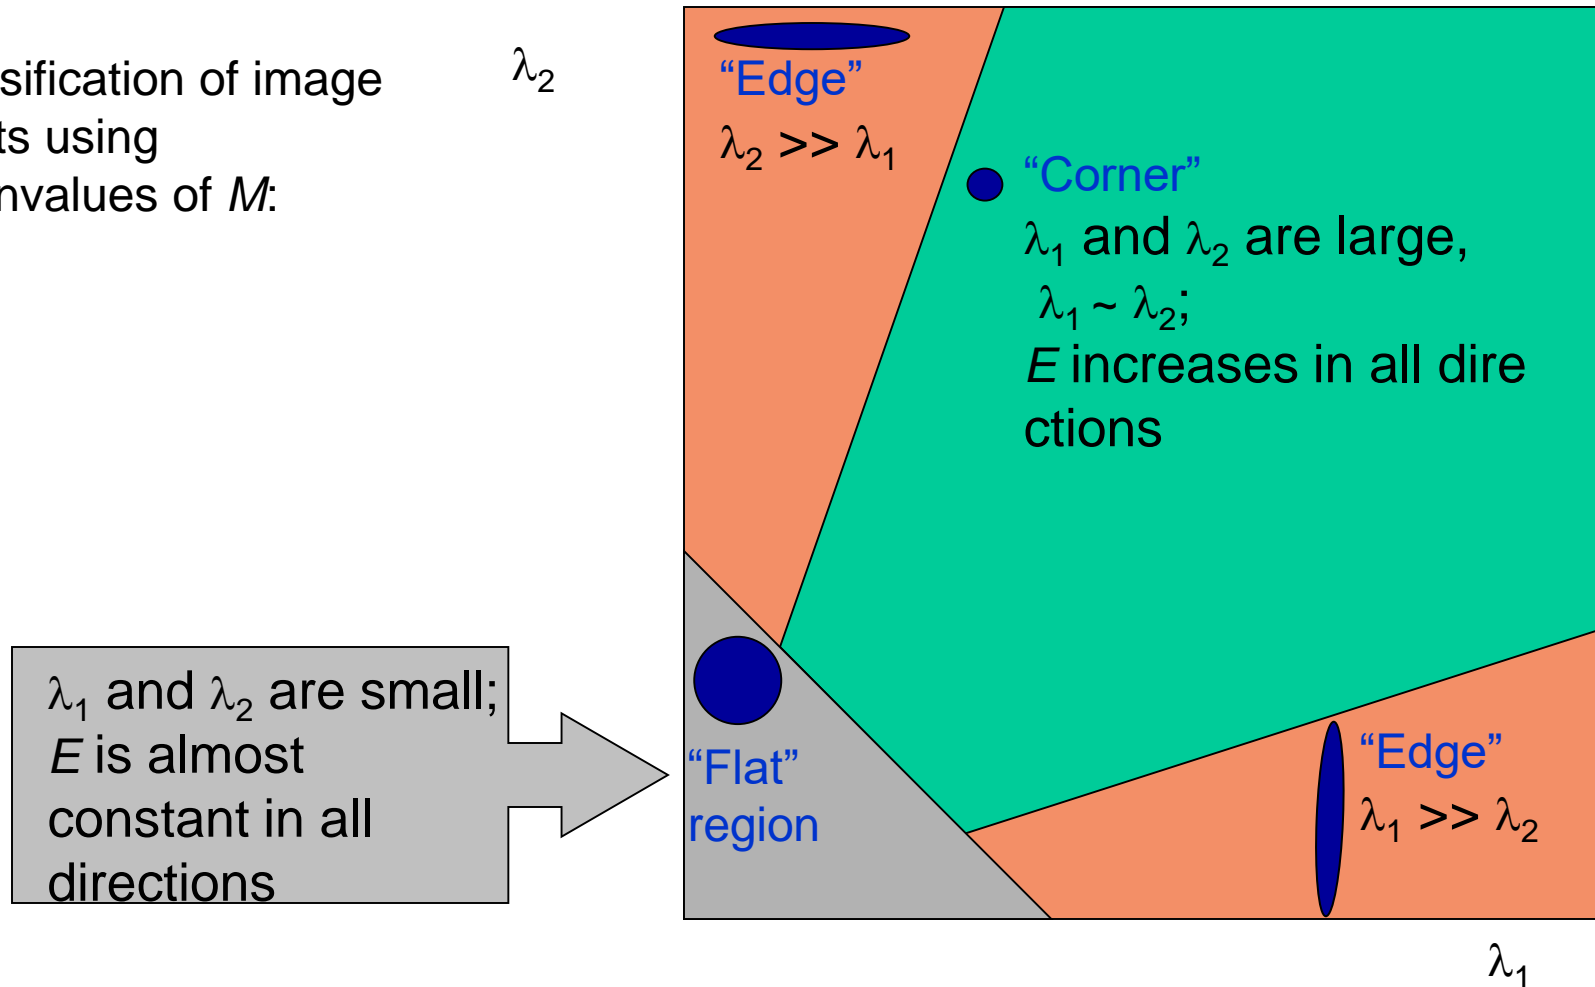
Selecting Good Features



small λ_1 , small λ_2

Harris Detector: Mathematics

Classification of image points using eigenvalues of M :



Harris Detector: Mathematics

Measure of corner response:

$$R = \det M - k (\text{trace } M)^2$$


This expression does not requires computing the eigenvalues.

$$\begin{aligned}\det M &= \lambda_1 \lambda_2 \\ \text{trace } M &= \lambda_1 + \lambda_2\end{aligned}$$

(k – empirical constant, $k = 0.04$ - 0.06)

Harris Detector: Mathematics

- R depends only on eigenvalues of M
- R is large for a **corner**
- R is negative with large magnitude for an **edge**
- $|R|$ is small for a **flat** region



Harris Detector


- The Algorithm:
 - Find points with large corner response function R ($R > \text{threshold}$)
 - Take the points of local maxima of R

Harris Detector: Workflow



Harris Detector: Workflow


✱
Compute corner response R



Harris Detector: Workflow




Find points with large corner response: $R > \text{threshold}$



Harris Detector: Workflow



Take only the points of local maxima of R



Harris Detector: Workflow



Harris Detector: Summary

- Average intensity change in direction $[u, v]$ can be expressed as a bilinear form:

$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

- Describe a point in terms of eigenvalues of M :
measure of corner response

$$R = \lambda_1 \lambda_2 - k (\lambda_1 + \lambda_2)^2$$

- A good (corner) point should have a *large intensity change in all directions*, i.e. R should be large positive

Ideal feature detector

- Would always find the same point on an object, regardless of changes to the image.
- i.e, insensitive to changes in:
 - Scale
 - Lighting
 - Perspective imaging
 - Partial occlusion



We want to:

detect *the same* interest points
regardless of *image changes*

SIFT Descriptor



- Basic idea:
 - Take 16x16 square window around detected feature
 - Compute edge orientation (angle of the gradient -90°) for each pixel
 - Throw out weak edges (threshold gradient magnitude)
 - Create histogram of surviving edge orientations






Image gradients



Keypoint descriptor

SIFT Descriptor

Full version

- Divide the 16x16 window into a 4x4 grid of cells (2x2 case shown below)
- Compute an orientation histogram for each cell
- $16 \text{ cells} * 8 \text{ orientations} = 128 \text{ dimensional descriptor}$





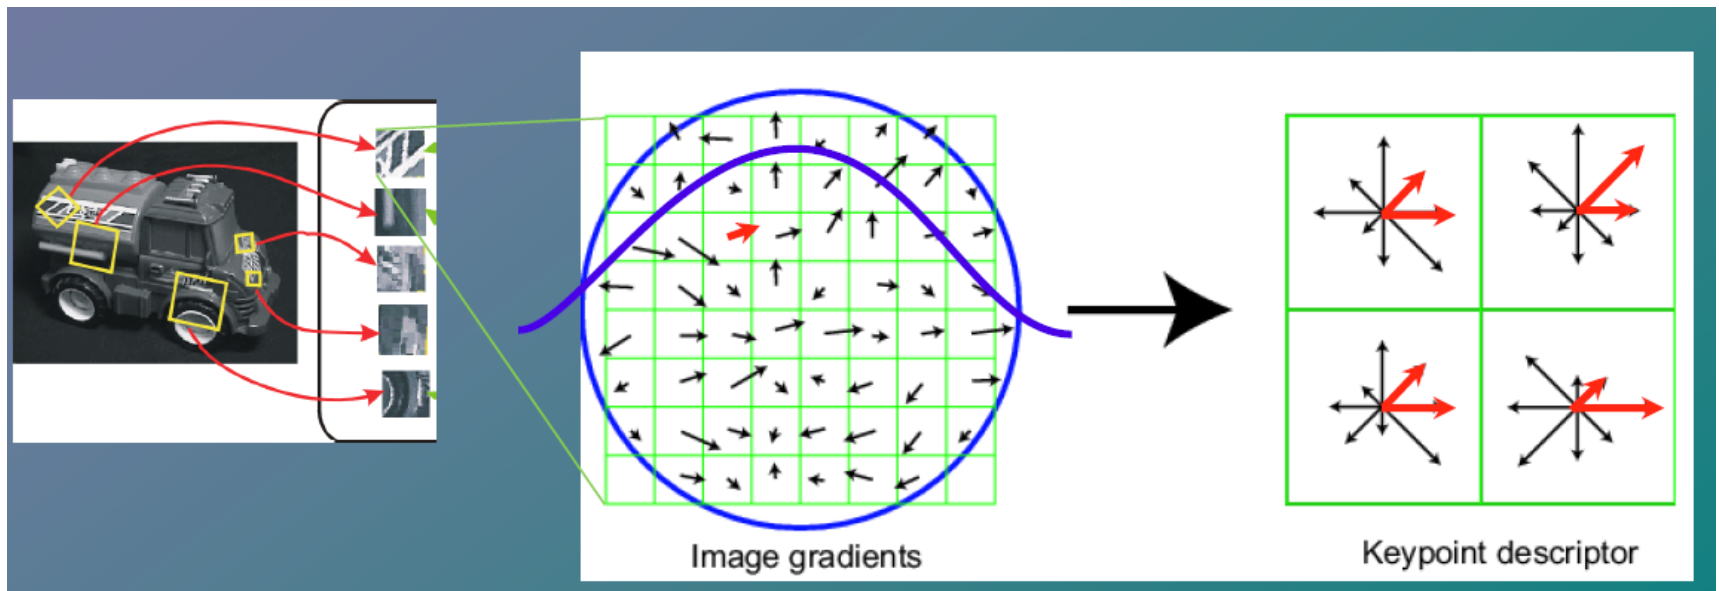
Image gradients



Keypoint descriptor
showing only 2x2 here but is 4x4

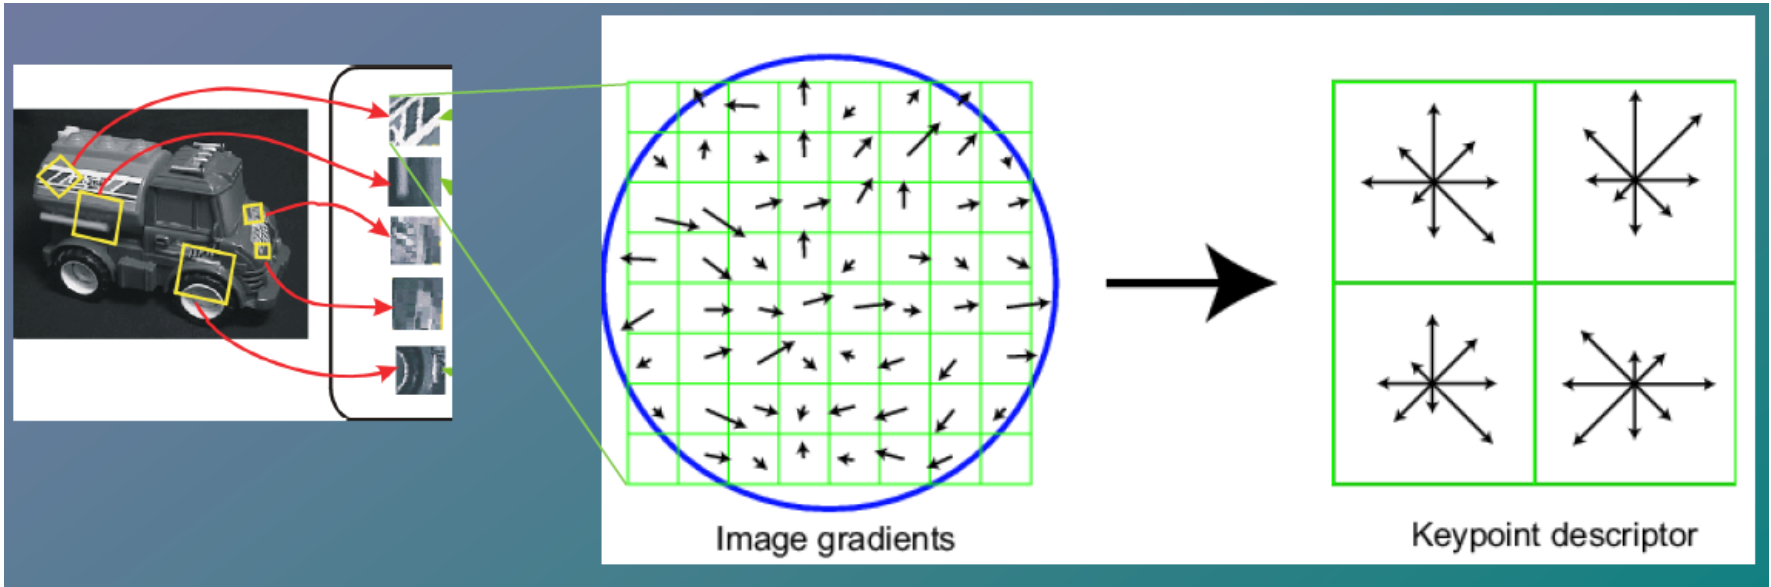
SIFT Descriptor

- Gaussian weight
- Trilinear interpolation
 - a given gradient contributes to 8 bins: 4 in space times 2 in orientation

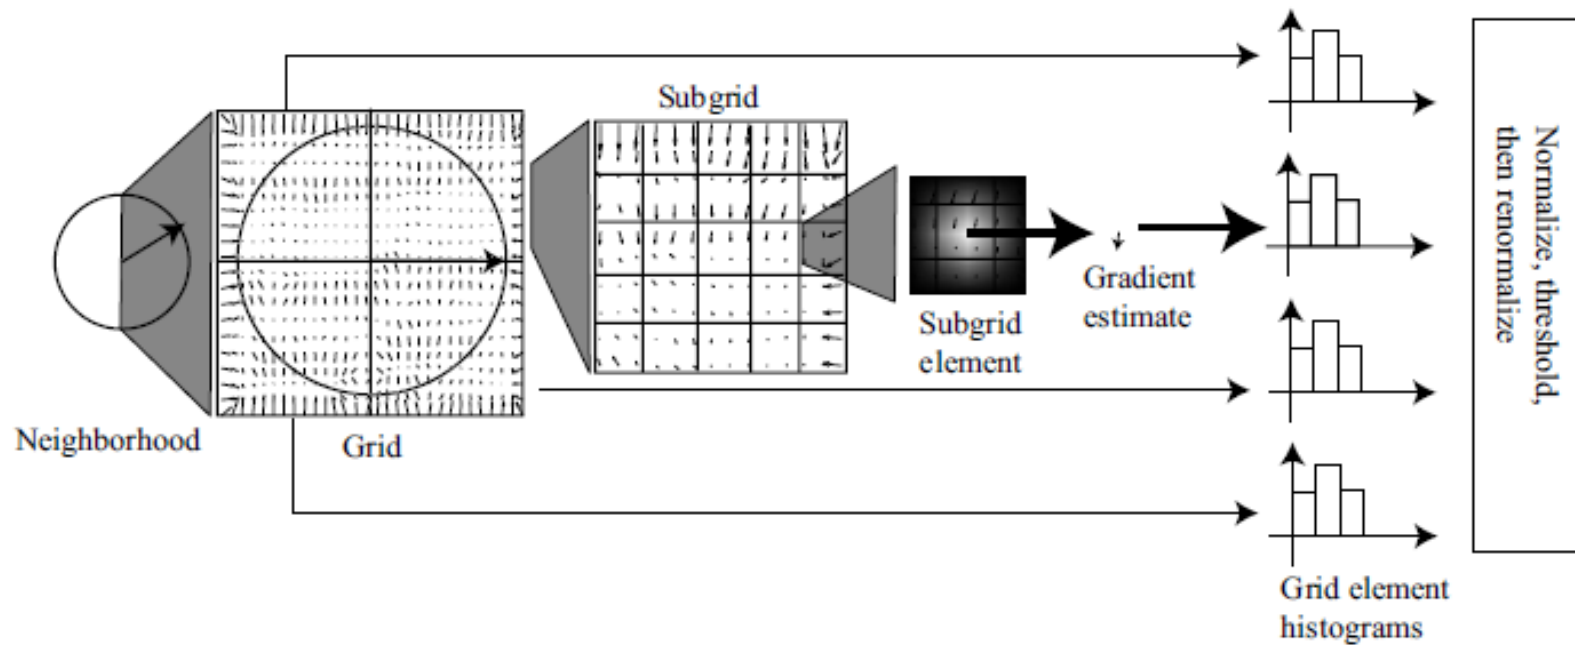


SIFT Descriptor

- 128-dim vector normalized to 1
- Threshold gradient magnitudes to avoid excessive influence of high gradients
 - after normalization, clamp gradients >0.2
 - renormalize



SIFT Descriptor




SIFT Descriptor

Given an image \mathcal{I} , and a patch with center (x_c, y_c) ,
radius r , orientation θ , and parameters n, m, q, k and t .
For each element of the $n \times n$ grid centered at (x_c, y_c) with spacing kr
Compute a weighted q element histogram of the averaged
gradient samples at each point of the $m \times m$ subgrid,
as in Algorithm 5.5.
Form an $n \times n \times q$ vector v by concatenating the histograms.
Compute $u = v / \sqrt{v \cdot v}$.
Form w whose i 'th element w_i is $\min(u_i, t)$.
The descriptor is $d = w / \sqrt{w \cdot w}$.


Example for scale invariance:

- Scale variation is just piece of cake!




Example for color invariance:

- Scale + color variation is still piece of cake!




Example for scale/color/geometrical invariance:

- Scale + color + geometrical variation is still piece of cake!



SIFT ORB + homography without RANSAC

- No corners need to be manually assigned
- ~~SIFT~~ ORB would lead to perform the transformation between the original image and the target image
- However, the transformation would fail every time. Do you know why?



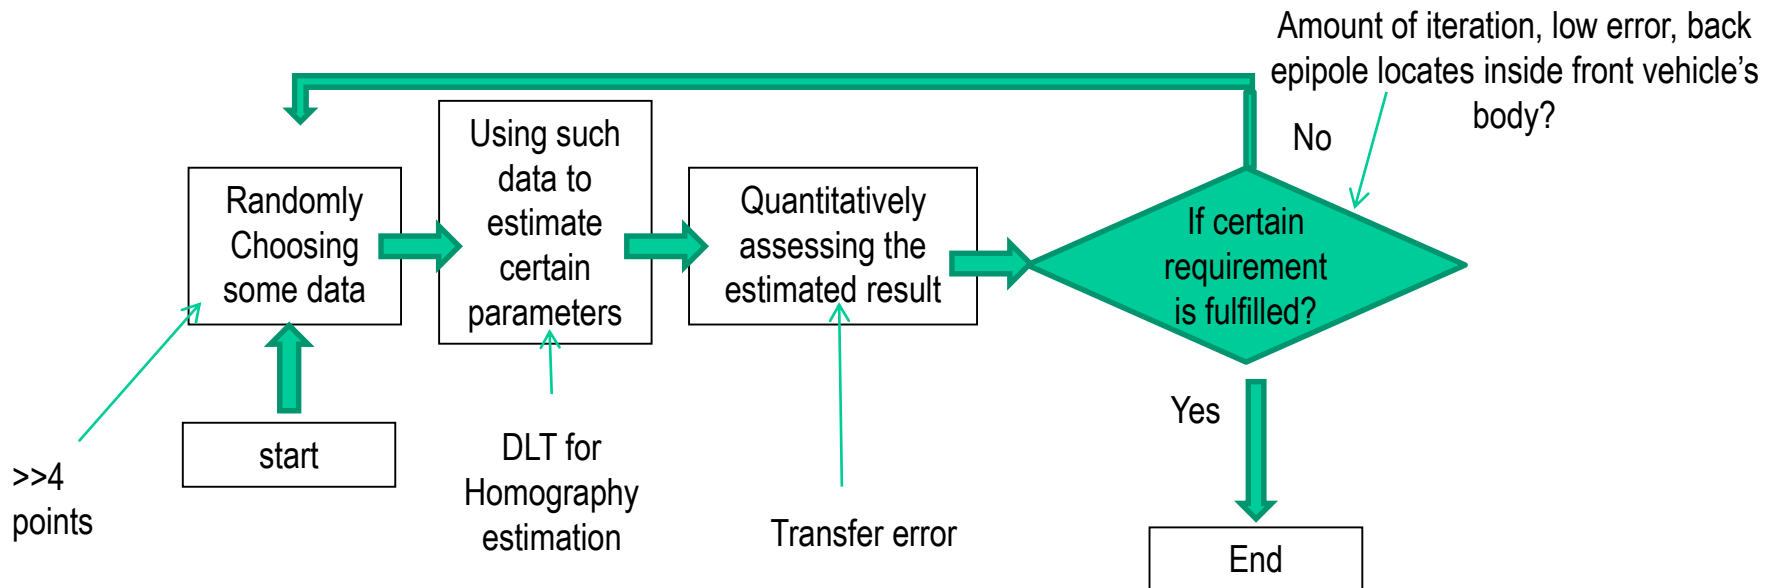
Before transformation



After transformation

How to obtain optimal H?

- The overall algorithm:
 - Using RANSAC+ Homography estimation
 - Input: feature correspondences,
 - Output: Optimal H




Cost function of homography estimation

- Algebraic error:
 - The standard SVD approach to minimize the norm $||Ah||$ under the constraint $||h||=1$
- Transfer error:
 - for a given correspondence (x_i, x'_i) , iterate to find H that minimizes the following cost function with LM optimization given the initial estimate from algebraic error
- $\sum_i d(x'_i, Hx_i)^2$
- Symmetric transfer error:
 - for a given correspondence (x_i, x'_i) , iterate to find H that minimizes the following cost function with LM optimization given the initial estimate from algebraic error
- $\sum_i d(x_i, H^{-1}x'_i)^2 + d(x'_i, Hx_i)^2$


Cost function of homography estimation (contd)

Symmetric transfer error:



SIFT ORB + homography with RANSAC = object identification

- The number of RANSAC is important but how large should we set?
- Giving more iterations is simply the answer if your application is not necessary to function in real-time.



Before transformation





After transformation


Panorama example: SIFT ORB + homography with RANSAC to produce Panorama



Panorama example: SIFT ORB + homography with RANSAC to produce Panorama



Panorama example: SIFT ORB + homography with RANSAC to produce Panorama



Panorama implementation review

- Could it be more efficient?
- Is it possible to control the seam line?
- Where is the best seam line?
- What if there is a moving object inside?
- Is it possible to alleviate parallax?
- What if the brightness of each image is different to some extent?
- What if the image is captured by a wide angle camera?



Without color blending



With color blending